

# Path Dependent Learning Dynamics of ICM PPO with Epstein Zin Utility in Time Series Reinforcement Learning

GitHub: [https://github.com/utilityfog/EZ\\_Optimization](https://github.com/utilityfog/EZ_Optimization)

# Define the Problem

## Big picture

Financial decision making over time is usually framed as predicting returns and then choosing an action based on those predictions. Classical models treat each decision in isolation and optimize a simple expected utility of wealth. Modern reinforcement learning (RL) instead treats the investor as an agent that repeatedly interacts with a stochastic environment and learns a policy that maximizes long run reward.

In this project we study a curiosity driven RL agent that allocates wealth between a risky asset and cash using a single asset time series. The agent uses Proximal Policy Optimization (PPO) together with an Intrinsic Curiosity Module (ICM). The value function is motivated by Epstein Zin recursive utility, which separates risk aversion from intertemporal substitution. Our question is not only whether this agent can earn good returns, but also how its learning path in early training episodes influences its long run behavior.

## Specific problem and hypothesis

We focus on the following questions.

1. Can an ICM PPO agent trained on a univariate financial time series learn a stable policy that produces positive risk adjusted returns out of sample
2. Are the learning dynamics path dependent in the sense that small stochastic differences during the first few episodes drive the agent either toward a good regime or into a degenerate regime
3. How do key state variables and hyperparameters relate to performance
  - fractionally differenced return feature
  - rolling volatility
  - curiosity bonus scale

Our working hypothesis is that the training dynamics of the agent are path dependent. In particular, we posit the existence of at least two basins of attraction in policy space. In a favorable regime, the stochastic trajectories observed in the first few episodes produce value

targets where intrinsic rewards and realized external returns are positively aligned, so gradient updates push the policy toward nontrivial trading behavior and stable Epstein Zin value estimates. In an unfavorable regime, early trajectories yield value targets dominated by noisy or misaligned intrinsic rewards, so the same architecture and hyperparameters drive the policy toward near zero allocation and effectively shut down exploration. Once the parameters enter one of these regions, subsequent PPO updates remain in the corresponding basin, so later episodes cannot reliably move the agent from the degenerate regime to the favorable regime.

## # Data Collection

### ## Source and structure

All raw data are loaded in `load_raw_data()` inside `src/data_preprocessing.py`. The sources are

- Daily S&P 500 OHLCV (open, high, low, close, volume) data stored in `sp500_df.csv`
- Monthly personal savings rate stored in `psavert_df.csv`
- Monthly unemployment rate stored in `unemploy_df.csv`

The preprocessing code converts each raw date column to a proper timestamp, then aggregates the daily S&P 500 data to month end by taking the last trading day in each calendar month. From the month end close prices it computes the monthly percentage return `SP500_Returns`. Savings and unemployment are already monthly series and are merged on the same month index.

The resulting merged dataset has the following columns for each month:

- Month end date
- S&P 500 open, high, low, close, and volume
- Personal savings rate
- Unemployment rate
- S&P 500 monthly return

After dropping the first row, whose return is undefined, we obtain a contiguous monthly panel. With the default chronological split of 80 percent for training and 20 percent for testing, the script reports

- 391 training months with 7 numeric features

- 98 test months with 7 numeric features

and stores

- features.npy and returns.npy for the train set
- features\_test.npy and returns\_test.npy for the test set

## ## Sampling limitations

The dataset reflects one historical period for a single broad equity index and two macro variables. There is strong temporal dependence, and the sample is not representative of all possible future regimes. Results therefore cannot be interpreted as general evidence that the strategy works on every asset or time period. The purpose of the project is to study learning dynamics in a controlled setting rather than to propose a ready to trade strategy.

## # Data Preparation

All data preparation steps are implemented in src/data\_preprocessing.py.

## ## Cleaning and filtering

- The first row is dropped because SP500\_Returns is NaN there.
- Any remaining rows with non finite values in SP500\_Returns are removed, though for this dataset there are no additional missing values after the first drop.
- The merged dataset is sorted chronologically so that all subsequent transformations respect time order.

## ## Feature construction

The feature matrix features\_raw consists of the following seven columns:

1. S&P 500 open price (month end)
2. S&P 500 high price
3. S&P 500 low price
4. S&P 500 close price
5. S&P 500 trading volume

6. Personal savings rate
7. Unemployment rate

All seven columns are cast to float32 and stacked into an array of shape (n\_months, 7).

The target array gross\_returns is defined as

$$\text{gross\_returns}_{t+1} = \text{SP500\_Returns}_t$$

so each element represents the gross portfolio return of investing in the S&P 500 between month t and t+1.

`## Expanding window normalization`

Rather than computing a single global mean and standard deviation, the script uses an expanding window z score transform implemented in `expanding_normalize()`.

For each feature dimension j and time index t:

- It computes the cumulative mean of feature j over months 0 through t.
- It computes the cumulative variance using cumulative sums of values and squared values.
- It defines the standard deviation as the square root of the variance, with a small floor of  $10^{-8}$  to avoid division by zero.
- It normalizes the current value by

$$z_{t,j} = \frac{x_{t,j} - \mu_{t,j}}{\sigma_{t,j}}$$

This produces an output matrix of the same shape where each row is normalized using only past and present information. From a time series perspective this is important, because it prevents the agent from peeking at future statistics during training.

`## Train test split`

After normalization, the script performs a chronological split:

- The first  $\lfloor 0.8 n \rfloor = 391$  months form the training set.

- The remaining 98 months form the test set.

Training features and returns are saved as `features.npy` and `returns.npy`. Test arrays are saved as `features_test.npy` and `returns_test.npy`. The RL environment only sees the training arrays during learning. The test arrays are reserved for out-of-sample evaluation.

There are no categorical variables, so no encoding is required.


## # Data Exploration

We explore the processed training dataset before training the RL agent. All visualizations mentioned here are generated in the accompanying code.

### ## Univariate patterns

#### 1. S&P 500 monthly gross return

Visualization: histogram of `gross_returns` - 1.



Observation: The histogram of S&P 500 monthly excess returns is centered slightly above zero, with most months between about  $-5\%$  and  $+5\%$ , a small positive mean around  $0.8\%$  per month, and a few extreme negative and positive outliers. This suggests a mildly right-shifted but heavy-tailed distribution, consistent with the idea that typical months are modestly positive while occasional shocks dominate long run performance.

#### 2. Covariance Matrix of features and returns

Visualization: table of Pearson correlations between features and returns

...

===== SUMMARY STATISTICS =====

	open	high	low	SP500_Close	volume	Personal_Savings_Rate
Unemployment	SP500_Returns					
count	489.000000	489.000000	489.000000	489.000000	4.890000e+02	
363.000000	363.000000	363.000000	363.000000	363.000000		
mean	1617.766602	1627.014526	1607.516113	1617.237183	2.387948e+09	
6.684023	8589.178711	0.008431				
std	1419.480103	1426.659668	1409.792480	1418.759399	2.034217e+09	
1.799131	2429.459961	0.043631				
min	179.539993	180.630005	178.860001	179.830002	1.499000e+07	2.200000
5481.000000	-0.217630					
25%	513.640015	515.289978	510.899994	514.710022	3.172200e+08	5.400000
6932.500000	-0.017004					
50%	1211.369995	1222.739990	1204.520020	1218.890015	1.799770e+09	
6.800000	7946.000000	0.012209				
75%	2073.169922	2075.760010	2057.939941	2065.300049	4.167160e+09	
8.100000	9135.000000	0.035794				
max	6860.500000	6880.750000	6820.689941	6822.339844	8.926480e+09	
12.000000	15352.000000	0.131767				

===== COVARIANCE MATRIX =====


	open	high	low	SP500_Close	volume
Personal_Savings_Rate	Unemployment	SP500_Returns			
open	2.014924e+06	2.025040e+06	2.001089e+06	2.013699e+06	2.025135e+12
-5.046762e+02	2.601919e+05	2.938607e+00			
high	2.025040e+06	2.035358e+06	2.011179e+06	2.024000e+06	2.037605e+12
-5.086552e+02	2.600097e+05	2.886005e+00			
low	2.001089e+06	2.011179e+06	1.987515e+06	2.000047e+06	2.008718e+12
-5.010075e+02	2.583550e+05	3.070056e+00			
SP500_Close		2.013699e+06	2.024000e+06	2.000047e+06	2.012878e+06
2.024120e+12		-5.040928e+02	2.586988e+05	3.026843e+00	
volume		2.025135e+12	2.037605e+12	2.008718e+12	2.024120e+12
-1.048551e+09		2.809062e+12	-5.477668e+06		4.138041e+18
Personal_Savings_Rate	-5.046762e+02	-5.086552e+02	-5.010075e+02	-5.040928e+02	
-1.048551e+09		3.236871e+00	7.837916e+02	6.771722e-03	
Unemployment		2.601919e+05	2.600097e+05	2.583550e+05	2.586988e+05
2.809062e+12		7.837916e+02	5.902276e+06	2.757018e+00	
SP500_Returns		2.938607e+00	2.886005e+00	3.070056e+00	3.026843e+00
-5.477668e+06		6.771722e-03	2.757018e+00	1.903655e-03	
...					

Observation: The covariance matrix shows that the four price levels (open, high, low, close) and trading volume move almost perfectly together, with covariances on the order of  $10^6$  between prices and  $10^{12}$  between prices and volume, indicating strong comovement and multicollinearity among these level variables. In contrast, the personal savings rate and unemployment have


much smaller variances and covariances that reflect slow-moving macro trends. Covariances involving SP500\_RetURNS are close to zero relative to its own variance (about  $1.9 \times 10^{-3}$ ), which suggests that one-month returns are only weakly linearly related to price levels, volume, or the macro variables. This supports treating return prediction as a difficult problem where simple linear structure is limited.

### 3. Training Wealth Trajectory and Policy Decisions (illustrates path dependence)


Visualization: Wealth Trajectory in a bad-seed round in the final episode




Visualization: Policy actions in a bad seed round in the final episode



Visualization: Wealth Trajectory in a good-seed round in the final episode



Visualization: Policy actions in a good seed round in the final episode



Observation: The wealth and policy plots exhibit a clear path-dependence phenomenon. In the bad-seed run, wealth falls by several orders of magnitude early in the episode and then hovers near the numerical floor, while the policy quickly drives the consumption fraction toward its upper bound, effectively liquidating wealth and preventing recovery. In the good-seed run, wealth grows steadily over the episode, ending far above the initial level, and the policy maintains moderate consumption rates that only rise after wealth has increased. With identical architecture and hyperparameters, the only difference between these runs is the stochastic training path, so the divergence in wealth and consumption policies provides visual evidence of multiple basins of attraction in the learning dynamics.

These explorations satisfy the requirement to examine at least three variables and their associations. They also motivated our final feature set: we retain the full 7 dimensional feature vector, but in this writeup we focus on the economic interpretation of returns, savings rate, unemployment, and volume.

## # Model Building

The key hyperparameters include:

- Curiosity weight that scales the intrinsic reward.
- Learning rate for the optimizer.
- Fractional differencing parameter used in feature construction.

These act as the three main “design variables” we vary in experiments, alongside the observed state features themselves.

```
src/config.py
...
class Config:
    # EZ parameters
    beta = 0.96
    psi = 1.5
    gamma_risk = 5.0

    # fracdiff preprocessing (still used by data_preprocessing)
    frac_d = 0.4
    frac_tol = 1e-4
    frac_max_lag = 512

    # learnable fractional differencing (inference-time)
    use_learnable_fracdiff = True
    fd_window = 12          # 12 months window in the state
    fracdiff_max_lag = 12    # kernel length in LearnableFracDiff1D
    fracdiff_init_d = 0.4

    # PPO
    gamma = 0.99
    gae_lambda = 0.95
    clip_ratio = 0.1 # 0.2
    vf_coeff = 1.0 # 0.5
    ent_coeff = 0.1
    ppo_epochs = 10
    batch_size = 128
    lr = 1e-2 # Changing this to any other value, such as 1e-4, fucks performance.

    # training
    num_episodes = 10
    start_wealth = 10000.0
    k_terminal: float = 1000.0

    device = "mps"
...

## Why this model
```

Standard DATA110 models such as linear regression or decision trees operate on independent and identically distributed samples and output a one step prediction. Our problem is inherently

sequential. Actions taken today change both wealth and the state distribution tomorrow. PPO is a natural choice because:

- It can handle continuous actions in bounded ranges.
- It uses a clipped objective that stabilizes policy updates.
- It is widely used in practice and relatively easy to implement.

The ICM adds an exploration pressure that is important when external rewards are sparse or noisy, which is very common in financial decision problems.

## 0) PURPOSE & SCOPE

This document replaces CRRA with **Epstein–Zin (EZ)** recursive preferences and **adds a Learnable Fractional Differencing (FracDiff) layer** in the feature pipeline, while preserving the **PPO + ICM** training stack and environment mechanics. It is **drop-in**: policy parameterization, buffers, rollout loop, exact log-probabilities, and PPO machinery remain intact. Only the **preference aggregator** (reward/value semantics) and **feature memory module** (FracDiff) change.

**What stays the same (do not touch):**

- Time/indexing, assets, returns, risk-free, ~~turnover/transaction cost~~, budget identity.
- Actor heads (consumption squashed Gaussian; ~~risky weights Dirichlet/softmax~~), exact log-prob math. (No more portfolio optimization)
- Critic backbone mechanics (but we output two EZ heads; see §5).
- ICM encoder/forward (and optional inverse) and curiosity reward wiring.
- PPO: ratio, clipping, GAE, epochs/minibatching, entropy, optimizers.
- Data split, standardization, rollout collection, buffer contents.

**What changes:**

1. **Utility/Value:** CRRA is replaced by **Epstein–Zin** with a numerically stable target in (z)-space (§4–§6).
2. **Features:** Insert **Learnable FracDiff** over returns before feature construction (§3).

## 1) Core definitions (time, single risky asset, wealth, consumption)

We now consider a **single risky asset** (S&P) and remove portfolio optimization entirely.

### 1.1 Time and assets

- Discrete time ( $t = 0, 1, 2, \dots, T - 1$ ).
- One risky asset with **gross return** ( $R_{t+1} \in R_{>0}$ ) between ( $t$ ) and ( $t + 1$ ).
- No explicit risk-free asset and no portfolio weights – all *unconsumed* wealth is automatically invested in the risky asset.

### 1.2 Wealth, consumption, normalization

- Wealth at start of step ( $t$ ): ( $W_t > 0$ ).
- Consumption fraction (action): ( $c_t \in (0, 1)$ ); **dollar consumption** ( $C_t := c_t \cdot W_t$ ).
- Running max wealth ( $M_t := \max_{0 \leq \tau \leq t} W_\tau$ ); normalized wealth ( $\tilde{W}_t := W_t / M_t \in (0, 1]$ ).

### 1.3 Budget identity (wealth transition)

In the simplified world, after consuming ( $C_t = c_t W_t$ ), the remaining wealth ( $(1 - c_t)W_t$ ) is fully invested in the risky asset, which realizes a gross return ( $R_{t+1}$ ) over ( $[t, t + 1]$ ).

The **wealth transition** is

$$([W_{t+1} = (1 - c_t)W_t R_{t+1} .])$$

We may optionally clip ( $W_{t+1}$ ) below by a small floor ( $\varepsilon_W > 0$ ) for numerical stability. Running max wealth is updated as

$$([M_{t+1} := \max(M_t, W_{t+1}) .])$$


---

## 2) OBSERVATIONS, FEATURES, STATE (BASELINE PIPELINE)

### 2.1 Observables at time ( $t$ )

- ( $W_t$ ) and a causal feature vector ( $x_t \in R^d$ ) built **only** from data ( $\leq t$ ).
- Standardize ( $x_t$ ) via train-set ( $(\mu, \sigma)$ ) to ( $\tilde{x}_t$ ) (store ( $\mu, \sigma$ ) from training only).

### 2.2 State to networks

- **State:** ( $s_t := \text{concat}(\tilde{W}_t, \tilde{x}_t) \in R^{1+d+n}$ ) (fixed order).

At time ( $t$ ) the agent observes:

- Normalized wealth ( $\tilde{W}_t = W_t / M_t$ ).
- A standardized feature vector ( $\tilde{x}_t \in R^d$ ), built from the FracDiff pipeline and other signals, using only information up to time ( $t$ ).

The **state** fed to the policy and critic is

$$(s_t := \text{concat}(\tilde{W}_t, \tilde{x}_t) \in R^{1+d}).)$$

There is no ( $w_{t-1}$ ) term in the state any more, since there is no portfolio decision.

---

## 3) Learnable Fractional Differencing (returns-domain feature module)

### 3.1 Goal & parameter

Learn a memory depth ( $d_{\text{target}} \in [d_{\min}, d_{\max}]$ ) (e.g., ([0, 1])) that controls the fractional differencing of returns to **capture long memory** while promoting **stationarity**.

### 3.2 Placement in pipeline

- Input raw **log-returns** per asset: ( $r_t \in R^n$ ) (or windows).
- Apply a FracDiff operator with effective exponent ( $d_{\text{eff}}$ ):
  - **Mode "direct"**: apply  $((1 - L)^{d_{\text{target}}})$  to returns.
  - **Mode "price\_equiv"**: apply  $((1 - L)^{d_{\text{target}} - 1})$  to returns (equivalent to price fracdiff of ( $d_{\text{target}}$ ) without reconstructing prices).

- Truncate the kernel to length ( $K$ ) (auto-chosen from  $(d_{\text{eff}})$  and a tolerance). Outputs lose the first ( $K$ ) steps.

### 3.3 State augmentation & alignment

- Build usual statistics **from** the FD output (lags, MAs, vol, PCA, cross-sectional transforms).
- **Shift** all time-aligned targets by ( $K$ ) (drop first ( $K$ ) steps) so shapes match.
- Optionally append ( $\text{stop}_{\text{grad}}(d_{\text{target}})$ ) and ( $K$ ) as scalar features so the policy/critic can adapt to memory depth.

### 3.4 Regularization & constraints

- Keep ( $d_{\text{target}}$ ) within bounds via a sigmoid reparameterization.
- Add a small L2 penalty if ( $d_{\text{target}}$ ) sticks to the bounds.
- Optional "whiteness" regularizer: penalize low-lag autocorrelation of FD residuals to avoid over-memory.

**Everything backpropagates end-to-end** because kernel weights are differentiable functions of ( $d_{\text{eff}}$ ).

## 4) Epstein–Zin Preferences (replace CRRA)

Let ( $\beta \in (0, 1)$ ) be the subjective discount, ( $\gamma > 0$ ) risk aversion, ( $\psi > 0$ ) intertemporal elasticity (EIS).

Define transforms:

- $(z(V) := V^{1 - \frac{1}{\psi}})$  (EIS/consumption space)
- $(y(V) := V^{1 - \gamma})$  (risk space)

### 4.1 EZ aggregator (Kreps–Porteus form)

For lifetime utility ( $V_t$ ) and consumption ( $C_t$ ):

$$(V_t = \left[ (1 - \beta)C_t^{1 - \frac{1}{\psi}} + \beta \left( E_t[V_{t+1}^{1 - \gamma}] \right)^{\frac{1 - \frac{1}{\psi}}{1 - \gamma}} \right]^{\frac{1}{1 - \frac{1}{\psi}}}.)$$

### 4.2 Practical RL parameterization (stable targets)

We train in ( $z$ )-space with a two-head critic predicting ( $\hat{z}_t \approx z(V_t)$ ) and ( $\hat{y}_t \approx y(V_t)$ ).

- **External (shaped) reward:**  $(r_t^{\text{ext}} := (1 - \beta)C_t^{1 - \frac{1}{\psi}}).$
- **One-step bootstrap target for ( $z$ ):**

$$(T_t^{(z)} := (1 - \beta)C_t^{1 - \frac{1}{\psi}} + \beta \left( \hat{y}_{t+1} \right)^{\frac{1 - \frac{1}{\psi}}{1 - \gamma}}.)$$

- **Value loss:**  $(L_{\text{value}} := \frac{1}{2} (\hat{z}_t - T_t^{(z)})^2).$

- Optional **consistency** regularizer: encourage  $(\hat{y}_t \approx (\hat{z}_t)^{\frac{1-\gamma}{1-\psi}})$  with a small weight.

**Degeneracies:**  $(\psi \rightarrow 1)$  approaches additive/separable (log-like);  $(\gamma \rightarrow 1)$  reduces risk curvature; recipe reduces toward CRRA smoothly.

## 5) ACTOR & CRITIC (Z-functions, distributions, exact log-probs)

We now have a **single action dimension**: the consumption rate ( $c_t \in (0, 1)$ ).

### Dimensions and symbols used throughout this section

- State at time ( $t$ ): ( $s_t \in R^{1+d+n}$ ) is the concatenation ( $s_t := \text{concat}(\tilde{W}_t, \tilde{x}_t)$ ), where ( $\tilde{W}_t = W_t / M_t$ ) and ( $\tilde{x}_t$ ) is the standardized feature vector.
- Consumption fraction (action component): ( $c_t \in (0, 1)$ ). Dollar consumption: ( $C_t := c_t W_t$ ).
- Hyperparameters for heads: ( $\sigma_{\min} > 0$ ) (std floor).

### 5.1 Actor ( $f_\theta$ )

The actor takes ( $s_t \in R^{1+d}$ ) and passes it through a shared backbone (e.g. an MLP) to produce parameters for a scalar Gaussian in a latent space:

- Pre-squash Normal parameters: ( $[\mu_c(s_t) \in R, \quad \ell_c(s_t) \in R, \quad \sigma_c(s_t) := \text{softplus}(\ell_c) + \sigma_{\min} .]$ )
- Sample pre-squash variable ( $[y_c \sim N(\mu_c(s_t), \sigma_c(s_t)^2) .]$ )
- Squash to the action space via the sigmoid: ( $[c_t := \sigma(y_c) = \frac{1}{1 + e^{-y_c}} \in (0, 1) .]$ )
- Deterministic (evaluation) action is given by ( $[c_t^{\text{det}} := \sigma(\mu_c(s_t)) .]$ )

There is no risky-weights head any more; ( $w_t$ ) is implicitly equal to (1) on the single asset.

### 5.2 Exact log-probability of ( $c_t$ )

Let

$$([y_c = \text{logit}(c_t) = \log \frac{c_t}{1 - c_t} .])$$

The log-probability under the squashed Gaussian is

$$([\log p(c_t | s_t) = \log N(y_c, \mu_c(s_t), \sigma_c(s_t)^2) - \log(c_t(1 - c_t)) .])$$

where the first term is the Gaussian log-density of ( $y_c$ ) and the second term is the log-Jacobian of the sigmoid.

This ( $\log p(c_t | s_t)$ ) is the **only** action log-probability used in PPO here.

## 5.3 Critic ( $g_\psi$ ) (two heads for EZ)

The critic takes ( $s_t$ ) and outputs two scalars:

- ( $\hat{z}_t \approx z(V_t)$ ) with ( $z(V) := V^{1-\frac{1}{\psi}}$ ).
  - ( $\hat{y}_t \approx y(V_t)$ ) with ( $y(V) := V^{1-\gamma}$ ). These are used to build the EZ bootstrap target and TD residual below.
- 

## 6) ENVIRONMENT STEP (FULL SEQUENCE)

At time ( $t$ ), given state ( $s_t$ ) and sampled consumption rate ( $c_t$ ), the environment performs:

### 1. Consumption and wealth evolution

Dollar consumption:

$$([C_t = c_t W_t .])$$

Remaining wealth:

$$([W_t^{\text{after}} = (1 - c_t) W_t .])$$

Apply the risky asset gross return ( $R_{t+1}$ ):

$$([W_{t+1} = W_t^{\text{after}} R_{t+1} = (1 - c_t) W_t R_{t+1} .])$$

Optionally clip ( $W_{t+1} \geq \varepsilon_W$ ) if needed for numerical stability.

### 2. Running max and next state

$$([M_{t+1} = \max(M_t, W_{t+1}), \quad \tilde{W}_{t+1} = \frac{W_{t+1}}{M_{t+1}} .])$$

The feature pipeline (including FracDiff) produces the next standardized feature vector ( $\tilde{x}_{t+1}$ ) from market data up to time ( $t+1$ ).

The next state is

$$([s_{t+1} = \text{concat}(\tilde{W}_{t+1}, \tilde{x}_{t+1}) .])$$

### 3. Termination

Episodes end when ( $t = T - 1$ ) (or when data runs out).

### Wealth transition

- Gross growth factor:

$$([G_{t+1} := (1 - c_t) (R_f[t+1] + w_t^\top \tilde{R}[t+1]) - \kappa \|w_t - w_{t-1}\|_1 .])$$

- Next wealth: ( $W_{t+1} := W_t \cdot G_{t+1}$ ). Safety floor may clip ( $G_{t+1} \geq \varepsilon_g > 0$ ).

## 7) REWARDS (EXTERNAL EZ FLOW, INTRINSIC ICM)

We use the Epstein–Zin flow term for external reward and the Intrinsic Curiosity Module (ICM) to supply an intrinsic shaping signal.

### EZ parameters

- Discount ( $\beta \in (0, 1)$ )
- Risk aversion ( $\gamma > 0$ )
- Elasticity of intertemporal substitution (EIS) ( $\psi > 0$ )
- Consumption ( $C_t = c_t W_t$ )

### 7.1 External reward (EZ flow term in $(z)$ -space)

The **external reward** at time ( $t$ ) is the EZ flow term in  $(z)$ –space, depending only on consumption:

$$([r_t^{\text{ext}} = (1 - \beta)C_t^{1 - \frac{1}{\psi}} = (1 - \beta)(c_t W_t)^{1 - \frac{1}{\psi}} .])$$

This is the main objective that encourages good consumption timing.

### 7.2 Intrinsic Curiosity Module (ICM) — complete specification

We define the ICM exactly and fully:

#### Network dimensions

Let:

- Feature dimension ( $d$ )
- State dimension ( $1 + d$ ) since state is ( $\text{concat}(\tilde{W}_t, \tilde{x}_t)$ )
- State-embedding dimension ( $m$ ) (e.g., 64)
- Hidden widths for ICM networks:
  - Encoder hidden width ( $E$ ) (e.g., 128)
  - Forward-model hidden width ( $F$ ) (e.g., 128)

Define the state encoder:

$$([\phi_\omega: R^{1+d} \rightarrow R^m .])$$

#### State encoder network

Given state ( $s_t \in R^{1+d}$ ):

$$([e1 = \text{GELU}(W_{e1}s_t + b_{e1}), \quad W_{e1} \in R^{E \times (1+d)} .])$$

$$([e2 = \text{GELU}(W_{e2}e1 + b_{e2}), \quad W_{e2} \in R^{E \times E} .])$$

$$([\phi(s_t) = W_{eo}e2 + b_{eo}, \quad W_{eo} \in R^{m \times E} .])$$

Define:

$$([\phi_t := \phi(s_t), \quad \phi_{t+1} := \phi(s_{t+1}) .])$$


---

## Action embedding (scalar action)

Because the action is **only** consumption ( $c_t \in (0, 1)$ ), we embed it as:

$$([y_c = \text{logit}(c_t) = \log \frac{c_t}{1 - c_t} .])$$

Then define:

$$([\psi(a_t) := \psi(c_t) := y_c \in R^1 .])$$

Dimensions: action embedding is 1-dimensional.

---

## Forward dynamics model

Maps  $((\phi(s_t), \psi(a_t)))$  into a prediction of  $(\phi(s_{t+1}))$ .

Input dimension to forward model:

$$([m + 1.])$$

Forward model layers:

$$([u1 = \text{GELU} \left( W_{f1}, \text{concat}(\phi_t, \psi(c_t)) + b_{f1} \right), \quad W_{f1} \in R^{F \times (m+1)} .])$$

$$([u2 = \text{GELU}(W_{f2}u1 + b_{f2}), \quad W_{f2} \in R^{F \times F} .])$$

$$([\hat{\phi}_{t+1} := W_{fo}u2 + b_{fo}, \quad W_{fo} \in R^{m \times F} .])$$

There is **no inverse model** in the consumption-only version.

---

## Intrinsic reward

Given:

- Encoded next state ( $\phi_{t+1}$ )
- Predicted next state ( $\hat{\phi}_{t+1}$ )

The intrinsic reward is:

$$([r_t^{\text{int}} := \eta \left| \phi_{t+1} - \hat{\phi}_{t+1} \right|_2^2,])$$

with a small scale factor ( $\eta > 0$ ) (e.g.,  $(10^{-3})$  ).

---

## ICM losses

Forward loss:

$$([L_{\text{fwd}}(\omega) := \left| \phi_{t+1} - \hat{\phi}_{t+1} \right|_2^2,])$$

Inverse loss:

$$([L_{\text{inv}} := 0])$$

since we removed portfolio weights and do not reconstruct ( $w_t$ ).

The action is 1-dimensional and directly known, so inverse dynamics is unnecessary.

Total ICM loss:

$$([L_{\text{ICM}} = L_{\text{fwd}},])$$


---

## 7.3 Total reward used by PPO

$$(r_t := r_t^{\text{ext}} + r_t^{\text{int}}).$$


---

# 8) Advantages, EZ targets, and losses (consumption-only)

---

All variables used below are defined here or earlier sections.

---

## 8.1 EZ one-step target in ( $z$ )-space

We have two critic heads:

- $(\hat{z}_t \approx z(V_t) := V_t^{1-\frac{1}{\psi}})$
- $(\hat{y}_t \approx y(V_t) := V_t^{1-\gamma})$

The one-step EZ bootstrap target is:

$$([T_t^{(z)} = (1-\beta)C_t^{1-\frac{1}{\psi}} + \beta(\hat{y}_{t+1})^{\frac{1-\frac{1}{\psi}}{1-\gamma}},])$$

All terms are fully defined:

- $(C_t = c_t W_t)$  is consumption
  - $(\hat{y}_{t+1})$  comes from the critic applied to next state
  - exponents come from EZ preference structure
- 

## 8.2 Value loss (z-head)

$$(L_{\text{value}} = \frac{1}{2} (\hat{z}_t - T_t^{(z)})^2 .])$$


---

## 8.3 TD residual in $(z)$ -space and GAE

Define the **combined reward**:

$$(r_t = r_t^{\text{ext}} + r_t^{\text{int}}])$$

and the **EZ temporal-difference residual**:

$$([\delta_t^{\text{EZ}} := r_t + \beta (T_t^{(z)} - r_t^{\text{ext}}) - \hat{z}_t .])$$

This matches the structure of the general EZ TD residual while incorporating intrinsic reward.

### Generalized Advantage Estimation (GAE)

Let  $(\lambda \in [0, 1])$  be the GAE parameter.

Compute the advantages by backward recursion:

$$([\tilde{A}_t = \delta_t^{\text{EZ}} + (\beta\lambda), \delta_{t+1}^{\text{EZ}} + (\beta\lambda)^2, \delta_{t+2}^{\text{EZ}} + \dots])$$

Practical implementation uses backward iteration over a rollout.

We normalize  $(\tilde{A}_t)$  to mean 0 and variance 1 in each minibatch.

---

## 8.4 PPO clipped policy loss (consumption-only)

Let:

- $(\log\pi_{\theta_{\text{old}}}(c_t | s_t))$  be the stored behavior log-prob.
- $(\log\pi_{\theta}(c_t | s_t))$  be recomputed with the current actor.
- Importance ratio:

$$([r_t(\theta) := \exp (\log\pi_{\theta}(c_t | s_t) - \log\pi_{\theta_{\text{old}}}(c_t | s_t)) .])$$

- Clipping parameter  $(\varepsilon \in (0, 1))$ .

The PPO objective (to minimize) is:

$$([L_{\text{PPO}} = -E_t [\min (r_t(\theta)\tilde{A}_t, \text{clip}(r_t(\theta), 1 - \varepsilon, 1 + \varepsilon)\tilde{A}_t)] .])$$

## 8.5 Entropy term (Gaussian action only)

The actor samples

$$([y_c \sim N(\mu_c(s_t), \sigma_c(s_t)^2)])$$

before applying the sigmoid.

Entropy of a Normal:

$$([H_c = \frac{1}{2} \log(2\pi e \sigma_c^2) .])$$

We encourage exploration by adding the entropy term:

$$([L_{\text{ent}} := -H_c .])$$

There is no Dirichlet entropy here since we removed risky-weight allocations.

## 8.6 ICM loss

As defined in §7.2:

$$([L_{\text{ICM}} = L_{\text{fwd}} = \left\| \phi_{t+1} - \hat{\phi}_{t+1} \right\|_2^2 .])$$

The inverse loss is zero in consumption-only and can be enabled later if desired.

## 8.7 Final training loss

Define scalar weighting hyperparameters:

- ( $c_v > 0$ ): value loss weight
- ( $\beta_{\text{ent}} > 0$ ): entropy loss weight
- ( $c_{\text{icm}} > 0$ ): curiosity loss weight

The full objective is:

$$([L_{\text{total}} = L_{\text{PPO}} + c_v L_{\text{value}} + \beta_{\text{ent}} L_{\text{ent}} + c_{\text{icm}} L_{\text{ICM}} .])$$

# 9) TRAINING PROCEDURE (COLLECT → TARGETS → PPO)

## 9.1 Hyperparameters (additions/changes)

- **EZ**: choose ( $\gamma \in [5, 10]$ ), ( $\psi \in [0.5, 1.0, 1.5]$ ), ( $\beta \in [0.95, 0.999]$ ).
- **FracDiff**: ( $d_{\text{target}}$ ) init (0.3) – (0.5) within ([0, 1]), tolerance ( $10^{-4}$ ), ( $K_{\text{max}} \in [1024, 4096]$ ) (match horizon).

- **RL:** keep PPO ( $\lambda$ ), clip, epochs, minibatch, lrs same initially.

## 9.2 Rollout collection (unchanged mechanics)

- Collect tuples  $(s_t, a_t = (c_t, w_t), r_t, s_{t+1}, \log\pi_{\theta_{\text{old}}})$  where  $(r_t)$  includes EZ flow + curiosity.
- Align time by dropping first ( $K$ ) steps due to FracDiff.

## 9.3 Target building & PPO update

- For each step, compute  $(T_t^{(z)}, (\delta_t^{\text{EZ}}))$ , GAE, and  $(z)$ -value loss.
- Recompute current  $(\log\pi_{\theta})$  exactly (§5.2); perform clipped PPO with entropy and ICM losses.
- After epochs, set  $(\theta_{\text{old}} \leftarrow \theta)$ .

## 9.4 Evaluation (deterministic)

- Use  $(c_t := \sigma(\mu_c)), (w_t := \alpha / \sum_i \alpha_i)$ .
  - Recover EZ value via inverse transform for reporting:  $(\hat{V}_t = \hat{z}_t^{1/(1-\frac{1}{\psi})})$ .
  - Report PnL, CAGR, MDD, Calmar, turnover, and  $(\hat{V}_0)$ .
- 

# 10) DIAGNOSTICS, CHECKS, & ABALATIONS

- **EZ sanity:** as  $(\psi \rightarrow 1)$  or  $(\gamma \rightarrow 1)$ , curves and training behavior should smoothly approach separable/CRRA.
  - **Scale hygiene:** track  $(\hat{z}_t, \hat{y}_t)$  magnitudes; clamp/normalize if exploding.
  - **FracDiff:** monitor learned  $(d_{\text{target}})$  trajectory; inspect ACF/PACF of FD outputs; avoid non-stationary drift.
  - **Alignment:** verify all post-FD tensors drop the first ( $K$ ) steps; shapes of policy/value/ICM batches match.
  - **Ablations:** (i) turn off FracDiff (identity) to test EZ alone; (ii)  $(\psi)$  grid with fixed  $(\gamma)$ ; (iii) compare CRRA vs EZ at matched  $(\gamma)$  with  $(\psi \approx 1)$ .
- 

# 11) MINIMAL MIGRATION CHECKLIST

- Expose  $(\gamma, \psi, \beta)$  in config; leave PPO hypers unchanged initially.
  - Critic: switch to **two heads**  $((\hat{z}, \hat{y}))$ ; keep shared backbone.
  - Reward pipe: compute  $(r_t^{\text{ext}} = (1 - \beta)C^{1-1/\psi})$ ; add curiosity as before  $\rightarrow (r_t)$ .
  - Targets: build  $(T^{(z)})$  with next-state  $(\hat{y})$ ; compute  $(\delta^{\text{EZ}})$  and GAE in  $(z)$ -space.
  - Insert **Learnable FracDiff** before feature builder; shift by  $(K)$ .
  - Log  $(d_{\text{target}}, K)$ , ACF diagnostics, and  $(\hat{z}, \hat{y})$  summaries.
-

```
# Model Evaluation
```

```
## Hypotheses and experimental setup
```

We evaluate this set of hypotheses.

### 1. Performance hypothesis

We train the agent with a fixed set of hyperparameters and measure

- average external reward per episode,
- the sequence of realized gross returns under the learned policy,
- simple risk measures such as mean, standard deviation, and Sharpe ratio of per step rewards.

### 2. Path dependence hypothesis

We keep data, architecture, and hyperparameters fixed, and only let the random seed that controls network initialization and the order of environment interactions vary. For each seed, we record

- per episode average external reward,
- per episode average curiosity reward,
- terminal wealth,
- summary statistics of the consumption fraction  $c_t$ .

```
## Metrics
```

Since this is an RL problem rather than a supervised prediction task, classical accuracy does not apply. Instead we print:

...

```
(venv) james@MacBook-Pro-502 EZ_Optimization % python -m src.eval
Non-finite values detected in train split during cleaning
keeping 363 rows out of 391
Warning: no fully finite rows in test split.
Falling back to column-wise imputation instead of dropping everything.
non-finite feature counts per column (test): [ 0  0  0  0  0 98 98]
non-finite returns in test: 0
/Users/james/Desktop/GitHub/EZ_Optimization/src/train.py:68: RuntimeWarning: Mean of empty
slice
    col_means = np.nanmean(
wealth: [ 10000.      10086.27717377  10001.62868886  9854.72473591
  10387.37248522  10513.40313031  10462.42804041  10336.92568461
  9977.97571284  10402.11748941  11078.88449752  11578.0997685
 11605.48823183  12435.14442234  13091.63060983  12906.39579346
 13554.66160685  13745.89725453  12939.23576564  13860.40002845
 12676.17405194  13369.91935279  13657.0544277  13270.7070526
 15019.33216571  15573.88510977  15984.86141853  15801.81525975
 15897.15154576  16658.8390342  17462.17299272  18072.6122216
```

17635.84882459 13797.73843349 12620.10478097 13539.61016067  
14087.03249909 14676.10127167 14186.73809018 14320.42932773  
14365.89797861 14987.29491194 14906.17796054 14330.78386115  
14900.12178577 15286.98068883 14998.18151383 15218.45265848  
16300.69324797 15828.86795782 16158.18651299 16967.53238609  
17563.71207087 17424.51001361 18964.2985533 19258.54372537  
19132.48991185 18650.80217516 18959.29288266 19365.32058513  
18032.63429182 18186.59822321 18627.6962139 18126.83300567  
19794.28534547 19618.36703128 19515.87774311 17675.23532837  
16770.52407017 16658.1752358 17656.5529593 18094.90679508  
18846.14748553 20114.11741338 20560.68757217 20567.24310487  
21361.2196443 20338.15431521 21250.4882297 21668.01294992  
21253.1844527 21504.67873944 20560.52303304 22854.80539714  
22399.42964669 22614.20754752 22120.4749859 22737.45221525  
22759.34668441 22364.24847365 23244.78701315 22686.94649819  
22893.501029 22941.69887309 23635.93130748 23874.81365112  
24043.01116948 24295.0447285 24749.27079118 24120.19961699  
24668.12500213 24686.72945765 24555.19756105 25400.65511091  
25146.92948248 25634.57742712 25303.59423341 25558.91193946  
26389.5692003 25596.66971223 24425.68710088 24707.30614309  
25013.58180563 24343.42365201 25109.96927917 26054.04615858  
25353.75259377 25881.94298792 24859.46264311 25165.18552093  
25776.112244 26705.93717473 27435.76180988 28202.84575016  
29226.91108453 29848.7860795 30797.23241587 30787.33863754  
32021.80115863 31862.32274406 33170.20617488 33748.79003098  
34849.55062154 35091.15148024 35368.91765818 35843.93695106  
36663.05520199 36745.75784491 35064.67802019 35724.34801566  
37660.68690266 38643.6295901 41479.10998297 40587.04390364  
43075.68009836 43330.97056256 41484.42382849 43907.30415116  
46479.21325816 48498.80837444 52288.73725389 49283.8836365  
51903.44924943 50113.8870267 52348.25739233 53171.72465863  
53711.37427674 57495.24190557 60366.82233742 60914.67944268  
59767.82491867 62124.90107618 61403.23807656 52450.79514929  
55723.43682573 60197.64827572 63756.8318433 67351.07279063  
70113.03152103 67849.53336045 70481.62852435 73155.91581005  
71329.10674382 75212.07059415 72801.74578852 72346.36572126  
70280.67525015 74675.92430959 76099.30602931 80501.10796971  
76403.23868894 74866.83957147 82107.86778267 79579.21458761  
77835.15204453 79697.94228164 78395.4983712 83153.95182492  
78706.55449268 78316.91728948 72046.12023929 72338.08493833  
74843.55661601 67936.12002699 63574.23873357 68457.58811065  
68805.98462952 67083.33348401 66362.78132302 62108.30540472  
57032.54698247 58064.72495472 62429.72940082 62902.49946742  
61922.80287359 60636.84101217 62864.50732005 59003.45868398

58467.56131603 54231.21966566 49946.67083708 50190.42942602  
44668.21623716 48529.68089198 51299.20194392 48204.13969823  
46882.59177388 46085.37304652 46470.48622745 50236.59369543  
52793.51676407 53391.20529081 54257.31112332 55227.00593274  
54567.30047933 57566.34048502 57976.642848 60919.79602399  
61972.20919882 62728.76136692 61702.52537782 60666.4331788  
61399.42923846 62503.88348752 60360.53417681 60498.54163058  
61064.98223602 61920.69988103 64310.46257781 66397.79714566  
64718.4996557 65941.831761 64681.11399763 63380.40522374  
65278.71296405 65269.33284534 67616.89241318 66858.0203972  
67322.55394554 66128.13278973 68454.85227513 68389.59001603  
70131.18387774 70162.89117177 70941.33098978 71803.59732131  
69583.58099974 69589.53358686 69943.38389709 71431.30263448  
73186.03174132 75491.89946156 76734.9189657 77702.91364602  
78795.26530204 77073.81305792 77842.92904083 81212.72538557  
83856.05580954 82361.96836763 79727.79760336 80753.29671449  
83643.69816407 84883.41060881 81144.77361133 80444.53649486  
75524.19756018 72898.81364901 72464.29122695 75909.65644644  
76719.85748147 70124.76314891 69433.30681041 70279.66056048  
63898.80340031 53072.72534824 49100.23532344 49484.22627115  
44552.83953646 39655.05140685 43041.72880834 47084.38120037  
48812.07737891 48821.53272876 52441.1937125 54200.71818619  
55791.82980812 54689.21811745 57826.35840593 58701.71398524  
56531.1841286 58143.03751263 61561.57757858 62470.11369334  
57349.01769367 54258.85957318 57990.60655311 54948.95549123  
59759.73815956 61962.17728532 61820.20665849 65819.12741891  
67309.57325998 69460.48577562 69387.66950332 71364.82968328  
70401.26779194 69115.85181417 67631.56128123 63790.62492301  
59212.82506912 65591.34249439 65259.46992971 65816.24767298  
68684.65084503 71472.45169668 73711.77102469 73159.04903336  
68575.51188741 71287.9417885 72185.92378938 73612.48588888  
75396.4975877 73904.37240605 74114.68318627 74638.47161418  
78402.27340973 79269.36997372 82122.01291557 83607.16488823  
85342.99848614 84062.826622 88220.66065681 85459.44131268  
88001.73004317 91926.13848724 94504.52335387 96731.21676242  
93289.14625567 97311.40138781 97985.87960113 98593.37090975  
100666.71908319 102585.15773624 101038.08729997 104842.60735617  
103215.99298351 105610.64968402 108201.55402424 107748.23354943  
104403.53910347 110134.40441304 108218.38791732]

Deterministic evaluation on train split:

split: train

initial\_wealth: 10000.0

final\_wealth: 108218.3879173154

pnl: 9.82183879173154

```
cagr: 0.0819115834643418
max_drawdown: -0.5328291933320188
calmar: 0.15372953375942502
V0_hat: 7.40093871508181
n_steps: 363
wealth: [10000.    10193.01944908 10419.17364791 10711.76051764
10817.06416238 11424.74122113 10979.76691689 10684.57073022
10713.60929446 10945.10179725 10998.09136612 11394.2487124
11739.06323575 11789.46222658 10971.22307796 11167.15071083
10142.25401136 10940.28074156 11265.51187292 11467.42691882
11918.24031674 11134.27451968 11901.74982044 12057.98585539
11839.82552937 12043.23540001 12289.28756958 12707.68832867
13070.98609353 13049.69215997 11952.06409598 10456.62011789
11782.96836429 12316.50968576 12542.97793298 13234.09878084
14161.32730242 13605.79368033 13229.36536721 14652.11093954
15196.00283636 15026.75522881 15418.80973907 16073.14677025
16915.7688383 17008.55999825 17386.37170018 17781.86106081
18297.34581347 17426.93909475 18631.8868459 18476.59613036
19282.3958405 18268.40922777 17695.48495659 18328.49211808
16716.36123238 16717.23327965 15314.30771533 16709.6760277
16000.49943147 14506.10787443 15664.59972712 16506.60087018
15533.16690755 16492.36875683 16061.69576786 16624.66702548
16868.07448072 16909.92910833 18004.4472817 18565.06734502
18236.14280539 17347.67249712 16966.359074 18479.38823727
19296.69877436 19603.41417448 20617.29427673 21256.79538644
20372.17200938 21350.44675079 22090.63945846 22340.71201182
22850.83169315 23312.32467871 23081.58522618 24404.16526642
23794.27793156 24437.08641898 24089.02659264 22702.80788246
22529.67816392 23915.76555178 25102.1250501 25645.9752691
26134.92691574 27058.08319614]
```

Deterministic evaluation on test split:

```
split: test
initial_wealth: 10000.0
final_wealth: 27058.0831961391
pnl: 1.7058083196139102
cagr: 0.12962509011033352
max_drawdown: -0.24770199748896937
calmar: 0.5233106370735099
V0_hat: 0.9194045989322054
n_steps: 98
...  
## Overfitting, underfitting, and complexity
```

To study model complexity we vary

- the PPO learning rate,
- the curiosity weight,
- the number of PPO epochs per update.

If the model underfits, all seeds produce low reward and the consumption policy reacts weakly to the state. If the model overfits or becomes numerically unstable, we observe high variance value estimates, occasional NaNs, and poor out of sample behavior.

Empirically we find that

- Very large curiosity weight pushes the agent toward chasing intrinsic reward at the expense of external utility.
- Very small curiosity weight reduces exploration and makes it more likely that the agent stagnates.
- Large learning rates tend to create exploding gradients and numerical issues.

## ## Path dependence results

When we train multiple agents with different random seeds, the most striking pattern is qualitative path dependence.

- For some seeds, the agent's early episodes include sequences of months where market returns and curiosity bonuses are aligned. The critic receives coherent value targets, and PPO updates gradually increase nontrivial consumption and investment behavior. These runs show improving external reward and relatively stable critic outputs.
- For other seeds, early episodes are unlucky. Either the market returns are unfavorable, or the curiosity module generates large bonuses on transitions that do not coincide with good consumption events. In these runs the critic's value targets become noisy and occasionally large in magnitude. PPO updates then shrink the consumption fraction  $c_t$  toward very small values, which effectively freezes wealth dynamics and stops exploration. Once this happens, later episodes rarely recover, even though the same architecture and hyperparameters are used.

These observations support the gateway hypothesis: the first few episodes act as a gate that determines which basin of attraction the learning trajectory enters.

## ## Class imbalance and normalization

There is no class label, so class imbalance does not apply in a strict sense. However, there is an imbalance between many small return months and a few extreme shocks. Expanding normalization and clipping of advantages help prevent these rare events from dominating the gradient updates.

We experimented with removing normalization and found that training quickly became unstable, with value estimates diverging. With expanding normalization in place, the training dynamics are stable enough to study path dependence itself rather than numerical artifacts.

## # Model Deployment

### ## Possible deployment scenario

In a real application, a model of this type could be deployed as an automated decision rule or a decision support tool for consumption and investment. A hypothetical deployment pipeline would be

1. Continuously collect new monthly S&P 500 and macro data.
2. Update the feature pipeline and recompute expanding normalized features using a rolling window so that statistics remain current.
3. Periodically retrain or fine tune the RL agent on the most recent history.
4. At each decision date, feed the current state into the policy network and implement the recommended consumption and investment decisions subject to human risk constraints.

### ## Risks and limitations

Several limitations make immediate deployment inappropriate.

- The model is trained and evaluated on one historical sample of one index with no transaction costs.
- The strong path dependence we observe means performance is sensitive to initialization and early data segments.
- Curiosity driven exploration can suggest actions that are useful for learning but unacceptable from a risk management perspective.
- Automated financial decision systems can affect market liquidity and volatility, so any real world deployment would require careful regulation and oversight.

For the DATA110 project we restrict ourselves to offline simulations and do not connect the model to any live trading system.

## # Project Meetings

2025/11/10, Zoom

2025/11/17, Davis Library Room 301

2025/12/03, Morrison Lounge

2025/12/04, Facetime, full group, final review of results

~2025/12/04, individual coding session(s) recorded in GitHub